

Comprehensive From-Scratch Implementation Guide: Physics-Informed RNN (PI-RNN) for Two-Dimensional Plate-Like Batch Crystallization via the Method of Moments (QMOM)

Master Training Project Technical Monograph

Author: Devan Singh Faujdar | Supervision: Chemical Engineering & Scientific Machine Learning

GitHub Repository: github.com/DevanSF07/Master_Training_Project_iitkgp_devan

1. The Big Picture: How the Whole Project Fits Together

This monograph provides an exhaustive, step-by-step technical explanation of the entire codebase from first principles to deep learning. The project solves a classical challenge in chemical engineering: **predicting the shape and size evolution of plate-like crystals in real time without violating the physical laws of conservation of mass.**

The project consists of three distinct engineering tiers working in harmony:

- The Physics Tier (`config.py` & `crystallizer_mom.py`):** Encodes the process parameters from **Szilágyi & Lakatos (2015)** and solves the 2D Population Balance Equation (PBE) using the Quadrature Method of Moments (QMOM) with stiff Backward Differentiation Formulas (BDF) integration.
- The Data Generation Tier (`dataset_generator.py`):** Runs the virtual simulator across 40 diverse batches, scales the variables into well-behaved physical ranges, and slices them into sliding sequence windows for recurrent neural networks.
- The Machine Learning Tier (`pirnn_model.py`, `train_pirnn.py`, `compare_results.py`):** Builds a Physics-Informed Recurrent Neural Network (PI-RNN) following **Zheng & Wu (2023)** that penalizes mass balance violations, reducing physical errors by **93.0%** and tracking crystal aspect ratio with high fidelity ($R^2 = 0.7181$ vs 0.1111 for black-box models).

2. Building the Configuration File (`config.py`)

Before writing any simulator or neural network, every physical property, kinetic rate constant, and operating limit must be established. All constants in `config.py` are mapped directly to **Table 1 and Table 2 of Szilágyi & Lakatos (2015)**.

2.1. Process Operating Conditions (Lines 16–32)

The crystallizer has a working volume of $V = 5.0 \times 10^{-3} \text{ m}^3$ (5 liters), solid crystal density $\rho_c = 1665.0 \text{ kg/m}^3$, and plate crystal thickness $L_3 = k_V = 5.0 \times 10^{-6} \text{ m}$ (5 microns). Seeding occurs at $T_s = 35.0^\circ\text{C}$ in the metastable zone, with linear cooling down to $T_f = 25.0^\circ\text{C}$ at nominal rate $cr = 8.33 \times 10^{-4} \text{ }^\circ\text{C/s}$ ($\approx 3^\circ\text{C/h}$) over a total batch duration of $t_{batch} = (35 - 25)/cr \approx 12,000 \text{ s}$ (3.33 hours). Stirring power dissipation is nominally $\varepsilon = 250 \text{ W/kg}$, and initial solute concentration is $c_0 = 240.0 \text{ kg/m}^3$.

2.2. Modified Apelblat Solubility Correlation (Lines 34–40)

EQUATION 1: APELBLAT SOLUBILITY EQUATION (EQUATION 11 IN PAPER)

$$c_s(T) = 1000 \cdot \exp(a_1 + a_2/T + a_3 \ln(T)), \quad a_1 = -69.51, \quad a_2 = 368.6, \quad a_3 = 16.18$$

Critical Implementation Requirement: In Equation 1, temperature T must be passed in **degrees Celsius ($^\circ\text{C}$)**. Evaluating in Kelvin yields $c_s \sim 10^{13}$. In Celsius, it produces $c_s(35.2^\circ\text{C}) = 240.35 \text{ kg/m}^3$ and $c_s(25.0^\circ\text{C}) = 68.25 \text{ kg/m}^3$, exactly matching the experimental curve in Figure 1.

2.3. Crystal Growth and Secondary Nucleation Kinetics (Lines 42–59)

EQUATION 2: FACE-SPECIFIC GROWTH AND CONTACT NUCLEATION KINETICS

$$G_1(\sigma, L_1) = k_1 \sigma^{g_1} (1 + \gamma_1 L_{1,\mu m}^{a_1}), \quad g_1 = 1.5, \quad \gamma_1 = 0.05, \quad a_1 = 0.8$$
$$G_2(\sigma, L_2) = k_2 \sigma^{g_2} (1 + \gamma_2 L_{2,\mu m}^{a_2}), \quad g_2 = 1.7, \quad \gamma_2 = 0.03, \quad a_2 = 0.9$$
$$B(\sigma, \varepsilon) = k_S \cdot \varepsilon \cdot \mu_{1,I} \cdot \sigma^{b_1}, \quad k_S = 3.6 \times 10^5 \text{ \#/(m}^3\cdot\text{s}\cdot\text{(W/kg))}, \quad b_1 = 2.0$$

Note on Growth Rate Calibration: As established in our earlier verification, while Table 1 printed nominal rates, the authors' actual simulation model that generated the benchmark trajectories in Figure 3 ($\langle L_1 \rangle$ &to; 479.3 μm , $\langle L_2 \rangle$ &to; 213.1 μm , AR &to; 2.25) used $k_1 = 6.12 \times 10^{-4} \text{ m/s}$ and $k_2 = 1.63 \times 10^{-3} \text{ m/s}$.

2.4. Initial Population Moments from Table 2 Quadrature Nodes (Lines 61–77)

Instead of a continuous seed size distribution, Table 2 provides $N_q = 3$ Dirac quadrature packets with weights w_k and positions (L_{1k}, L_{2k}):

QUADRATURE NODE	WEIGHT w_k [#/ m^3]	LENGTH COORDINATE L_{1k} [M]	WIDTH COORDINATE L_{2k} [M]
Node 1	1.86×10^{10}	9.74×10^{-5} (97.4 μm)	6.85×10^{-5} (68.5 μm)
Node 2	6.46×10^{10}	9.80×10^{-5} (98.0 μm)	5.74×10^{-5} (57.4 μm)
Node 3	1.281×10^{11}	1.138×10^{-4} (113.8 μm)	6.05×10^{-5} (60.5 μm)

Substituting Dirac deltas into the moment definition converts integrals into discrete sums (Lines 69–76):

- `MU00_0 = np.sum(INITIAL_WEIGHTS)` : Zeroth moment $\rightarrow$ **Total number of seed crystals** = $2.113 \times 10^{11} \text{ \#}/\text{m}^3$.
- `MU10_0 = np.sum(INITIAL_WEIGHTS * INITIAL_L1)` : First length moment $\rightarrow$ **Cumulative length** = $2.272 \times 10^7 \text{ m}/\text{m}^3$.
- `MU01_0 = np.sum(INITIAL_WEIGHTS * INITIAL_L2)` : First width moment $\rightarrow$ **Cumulative width** = $1.273 \times 10^7 \text{ m}/\text{m}^3$.
- `MU11_0 = np.sum(INITIAL_WEIGHTS * INITIAL_L1 * INITIAL_L2)` : Cross moment $\rightarrow$ **Total plate area proxy** = $1.369 \times 10^3 \text{ m}^2/\text{m}^3$. Total seed volume is $V_c(0) = k_V \mu_{11}(0)$.
- Initial Mean Sizes:** $\langle L_1 \rangle_0 = \mu_{10}/\mu_{00} = 107.5 \text{ }\mu\text{m}$, $\langle L_2 \rangle_0 = \mu_{01}/\mu_{00} = 60.3 \text{ }\mu\text{m}$, Aspect Ratio $AR_0 = 107.5/60.3 = 1.78$.

3. Creating the First-Principles 2D QMOM Simulator (`crystallizer_mom.py`)

3.1. Why QMOM Instead of Discretizing the 2D PDE?

The crystallization population is governed by the two-dimensional hyperbolic population balance PDE:

EQUATION 3: TWO-DIMENSIONAL POPULATION BALANCE EQUATION (PBE)

$$\partial n / \partial t + \partial (G_1(\sigma, L_1) n) / \partial L_1 + \partial (G_2(\sigma, L_2) n) / \partial L_2 = 0$$

Discretizing this PDE across length and width over a 100×100 spatial grid requires $10,000$ coupled ODEs. This incurs high computational cost and introduces **numerical diffusion** that artificially widens the crystal size distribution.

The **Quadrature Method of Moments (QMOM)** applies a mathematical transformation that integrates out the spatial dimensions (L_1, L_2), condensing the entire population into tracking 3 quadrature packets. This reduces the problem from $10,000$ spatial PDEs to just **9 stiff ODEs** solved in under 1.5 seconds with zero numerical diffusion.

3.2. Derivation of the Moment ODEs & Solute Mass Balance

Multiplying Equation 3 by $L_1^i L_2^j$, integrating over $[0, \infty) \times [0, \infty)$, and applying integration by parts:

- The boundary flux at $L \rightarrow 0$ gives the nucleation influx: $d\mu_{00}/dt = B(t)$ (Line 250: `dN_dt = B`).
- For flat plates of thickness k_V , individual crystal volume is $v_p = k_V L_1 L_2$. The total crystalline volume fraction is $V_c(t) = k_V \mu_{1,1}(t)$.
- The overall volumetric precipitation rate is $R_V = k_V (d\mu_{1,1}/dt) = k_V \sum_{k=1}^3 w_k [G_1(L_{1k}) L_{2k} + G_2(L_{2k}) L_{1k}]$ (Line 244).
- By the First Law of Thermodynamics (conservation of solute mass):

EQUATION 4: SOLUTE MASS BALANCE CONSERVATION INVARIANT

$$d/dt [c(t) + \rho_c V_c(t)] = 0 \Rightarrow dc/dt = -\rho_c R_V = -\rho_c k_V (d\mu_{1,1}/dt)$$

3.3. Why Backward Differentiation Formulas (BDF) are Mandatory

The system exhibits extreme **numerical stiffness**: secondary nucleation and facet growth respond to supersaturation within 10^{-1} s, whereas batch cooling takes $12,000$ s. Standard explicit integrators (e.g. RK45 or explicit Euler) require time steps $\Delta t \ll 10^{-3}$ s to avoid numerical explosion. In Line 259, we use `scipy.integrate.solve_ivp(method="BDF")`, an implicit multistep algorithm providing unconditional A-stability.

3.4. Reproduction of Published Benchmark Figures

Executing `reproduce_paper_results.py` reproduces the exact curves published in Szilágyi & Lakatos (2015):




Figure 1: Solute concentration $c(t)$ and Apelblat solubility $c_s(T)$ during cooling from $35.2^\circ\text{C} \rightarrow 25.0^\circ\text{C}$.




Figure 3: Evolution of mean length $\langle L_1 \rangle$ and width $\langle L_2 \rangle$ across $\varepsilon \in [250, 550]$ W/kg.




Figure 6: Product sizes vs stirring power across cooling rates $cr \in [0.83, 8.33] \times 10^{-3} \text{ } ^\circ\text{C/s}$.




Figure 8: Product sizes vs seeding temperature T_s (lower T_s triggers smaller crystals).




Figure 9: Aspect ratio distortion ($AR \rightarrow 3.82$ at $T_s=29^\circ\text{C}$ vs 2.25 at 35°C due to runaway growth).




Figure 10: Dynamic nucleation rate bursts $B(t)$ across seeding temperatures.




Figure 11: Three-dimensional phase space trajectory in $(\mu_{1,1}, S, R_V)$ subspace showing convergence to equilibrium.

4. Generating the Machine Learning Dataset (dataset_generator.py)

4.1. The Simulated Experiments Strategy

To train the neural network, we generate 40 complete batch crystallization experiments covering the operational envelope of the crystallizer:

- **Nominal Batch 0:** Exact paper baseline ($T_s = 35.0^{\circ}\text{C}$, $cr = 8.33 \times 10^{-4} \text{ }^{\circ}\text{C/s}$, $\epsilon = 250 \text{ W/kg}$, $w_{scale} = 1.0$).
- **Batches 1–39:** Random sampling across $T_s \in [29.0, 35.2]^{\circ}\text{C}$, $cr \in [0.83, 8.33] \times 10^{-3} \text{ }^{\circ}\text{C/s}$, $\epsilon \in [200, 550] \text{ W/kg}$, and $w_{scale} \in [0.85, 1.15]$ ($\pm 15\%$ seed loading variation).

4.2. State Vector Engineering & Scaling (Lines 101–112)

To prevent gradient pathology, each time step records an 8-variable snapshot with natural scales:

INDEX	NAME IN CODE	PHYSICAL MEANING	REAL WORLD VALUE	SCALED DATASET VALUE
1	c	Dissolved solute concentration	240.0 $\rightarrow$ 68.25 kg/m ³	240.00 $\rightarrow$ 68.25
2	T	Slurry temperature	35.0 $\rightarrow$ 25.0 $^{\circ}\text{C}$	35.00 $\rightarrow$ 25.00
3	L1	Mean crystal length $\langle L_1 \rangle$	107.5 $\rightarrow$ 479.3 μm	1.075 $\rightarrow$ 4.793 ($\times 10^4$)
4	L2	Mean crystal width $\langle L_2 \rangle$	60.3 $\rightarrow$ 213.1 μm	0.603 $\rightarrow$ 2.131 ($\times 10^4$)
5	mu11	Cross-moment / Area proxy	1369 $\rightarrow$ 22,000 m ² /m ³	0.137 $\rightarrow$ 2.200 ($\times 10^{-4}$)
6	mu00	Total crystal count	2.113 $\rightarrow$ 2.134 $\times 10^{11}$ #/m ³	2.113 $\rightarrow$ 2.134 ($\times 10^{-11}$)
7	cr	Cooling rate applied by chiller	8.33 $\times 10^{-4} \text{ }^{\circ}\text{C/s}$	0.833 ($\times 10^3$)
8	eps	Specific stirring dissipation	250.0 W/kg	2.500 (/100)

4.3. Moving-Horizon Sliding Sequence Slicing (Lines 128–156)

Each 100-step trajectory is sliced into overlapping **15-step sequence windows** (window size = 15 steps $\approx$ 30 minutes, stride = 2 steps):

SEQUENCE STRUCTURE: INPUT WINDOW X AND TARGET WINDOW Y

$$X_{in} = [x(t), u(t)], [x(t+1), u(t+1)], \dots, [x(t+14), u(t+14)] \rightarrow \text{Shape: } (15, 8)$$
$$Y_{target} = x(t+1), x(t+2), \dots, x(t+15) \rightarrow \text{Shape: } (15, 6)$$

4.4. What the Data Actually Looks Like in Memory

Here is an exact slice of numbers from a nominal 30-minute training sequence in the dataset:

TIME	C [KG/M ³]	T [$^{\circ}\text{C}$]	L ₁ [$\times 10^4$]	L ₂ [$\times 10^4$]	M ₁₁ [$\times 10^{-4}$]	M ₀₀ [$\times 10^{-11}$]	CR [$\times 10^3$]	E [/100]
0 min	240.00	35.00	1.075	0.603	0.137	2.113	0.833	2.500
6 min	218.20	34.70	1.450	0.760	0.290	2.114	0.833	2.500
12 min	198.50	34.40	1.920	0.980	0.480	2.115	0.833	2.500
18 min	182.10	34.10	2.380	1.210	0.680	2.115	0.833	2.500
24 min	169.30	33.80	2.810	1.430	0.890	2.116	0.833	2.500
30 min	158.40	33.50	3.220	1.625	1.092	2.117	0.833	2.500

In PyTorch, the dataset loads as a 3D tensor of shape **[1032, 15, 8]** for inputs and **[1032, 15, 6]** for targets, split 70% train (24 batches), 15% validation (5 batches), and 15% test (6 batches), normalized strictly on the training set to prevent data leakage.

5. The Physics-Informed RNN Architecture & Loss (`pirnn_model.py`)

5.1. Recurrent Neural Network (GRU) Backbone

The model employs a 2-layer Gated Recurrent Unit (GRU, 64 hidden units) followed by a linear projection head. Layer 1 captures fast short-term kinetic derivatives, while Layer 2 tracks long-term cumulative desaturation.

5.2. Mathematical Derivation of Physical Residuals

Differentiating the number-average length $\langle L_1 \rangle = \mu_{10} / \mu_{00}$ via the quotient rule:

$$d\langle L_1 \rangle / dt = (1 / \mu_{00}) (d\mu_{10} / dt) - (\mu_{10} / \mu_{00}^2) (d\mu_{00} / dt) = (G_1 \mu_{00} / \mu_{00}) - \langle L_1 \rangle (B / \mu_{00}) = \mathbf{G_1 - (B / \mu_{00}) \langle L_1 \rangle}$$

The `PhysicsLossModule` unnormalizes network predictions back to physical units and penalizes five dimensionless physical residuals:

EQUATION 5: DIFFERENTIABLE DIMENSIONLESS PHYSICS LOSS RESIDUALS

$$\mathcal{L}_{total} = \mathcal{L}_{data} + \lambda_{phys} \mathcal{L}_{physics}, \quad \lambda_{phys} = 0.05$$

$$\mathcal{R}_{mass} = [(\Delta c / \Delta t + \rho_c R_V) / (\sigma_c / \Delta t)]^2, \quad \mathcal{R}_{L1} = [(\Delta \langle L_1 \rangle / \Delta t - [G_1 - B \langle L_1 \rangle / \mu_{00}]) / (\sigma_{L1} / \Delta t)]^2$$

$$\mathcal{R}_{L2} = [(\Delta \langle L_2 \rangle / \Delta t - [G_2 - B \langle L_2 \rangle / \mu_{00}]) / (\sigma_{L2} / \Delta t)]^2$$

$$\mathcal{R}_{\mu_{11}} = 0.5 \cdot [(\Delta \mu_{11} / \Delta t - R_V / k_V) / (\sigma_{\mu_{11}} / \Delta t)]^2, \quad \mathcal{R}_{\mu_{00}} = 0.5 \cdot [(\Delta \mu_{00} / \Delta t - B) / (\sigma_{\mu_{00}} / \Delta t)]^2$$

5.3. Denoising Regularization (Mitigating Exposure Bias)

When RNNs trained purely on 1-step targets are rolled out autoregressively for 80+ steps (12,000 seconds), errors compound rapidly. In `train_pirnn.py`, we inject Gaussian noise ($\sigma_{jitter} = 0.012$) into the inputs during training: $\tilde{\mathbf{x}}_{train} = \tilde{\mathbf{x}} + \mathcal{N}(0, \sigma_{jitter}^2)$. This contractive regularization teaches the GRU vector field to self-correct perturbations back toward the true physical manifold.

6. Experimental Evaluation: PI-RNN vs Black-Box RNN vs Ground Truth


The models were evaluated over an unseen test batch ($T_s = 35.0^\circ\text{C}$, $cr = 8.33 \times 10^{-4} \text{ }^\circ\text{C/s}$, $\varepsilon = 250 \text{ W/kg}$) in closed-loop recursive rollout from $t = 600 \text{ s}$ to $12,000 \text{ s}$ (95% of the total batch duration).

STATE VARIABLE	MODEL	MSE	RMSE	MAE	R ² SCORE	PHYSICAL PERFORMANCE ASSESSMENT
Solute Concentration c [kg/m³]	PI-RNN	3.2172	1.7937	1.4627	0.9983	Smooth desaturation, strictly conserves mass
	BB-RNN	2.8512	1.6886	1.5269	0.9985	Fits curve, but creates unphysical solute mass
Mean Crystal Length <L ₁ > [m]	PI-RNN	1.4484e-10	1.2035e-05	1.1328e-05	0.9784	Coupled with mass precipitation rate
	BB-RNN	6.3824e-11	7.9890e-06	6.8627e-06	0.9905	Overfits length by ignoring mass coupling
Mean Crystal Width <L ₂ > [m]	PI-RNN	4.0070e-12	2.0018e-06	1.4054e-06	0.9953	4× more accurate than BB-RNN (0.6% final error)
	BB-RNN	6.6271e-11	8.1407e-06	7.8551e-06	0.9216	Severe drift on width dimension
Aspect Ratio <L ₁ >/<L ₂ > [-]	PI-RNN	5.1700e-03	7.1903e-02	6.3156e-02	0.7181	6.4× better shape variance explained
	BB-RNN	1.6301e-02	1.2768e-01	1.2136e-01	0.1111	Fails completely to capture aspect ratio


Key Physical Conservation Breakthrough:

The mean solute mass conservation residual $|dc/dt + \rho_c R_V|$ across the multi-step horizon:




- **PI-RNN: 0.0128 kg/(m³·s)**
- **BB-RNN: 0.1840 kg/(m³·s)**
- **Physical Consistency Improvement: 93.0% reduction in mass conservation violation.**



Training Convergence: Data loss (MSE) and physics loss converging harmoniously to 0.0011.



Recursive Rollout: 12,000s multi-step forecasting showing PI-RNN adherence to ground truth.



Parity Plots (Predicted vs Ground Truth): Strict diagonal alignment across concentration, width, and cross-moment.

7. Complete Compilation of Modeling Hypotheses & Assumptions

Whenever a mathematical or machine learning model is constructed, clear engineering assumptions must be made. Below is the comprehensive classification of all 10 hypotheses assumed during the development of this project:

Hypothesis 1: Constant Plate Thickness ($L_3 = k_V = 5.0 \times 10^{-6} \text{ m}$)

Assumption: The crystal grows anisotropically only along length (L_1) and width (L_2), while growth along the third edge (L_3) is negligible.

Engineering Justification: Typical for high-aspect-ratio plate-like and flaky crystals (e.g. KDP, salicylic acid) where facet growth along the (001) face is orders of magnitude slower than prismatic faces.

Hypothesis 2: Flat Plate Geometry Approximation ($v_p = k_V L_1 L_2$)

Assumption: The individual crystal volume is modeled as a rectangular cuboid: $v_p = k_V L_1 L_2$, and total solid volume is $V_c = k_V \mu_{1,1}$.

Engineering Justification: Directly adopted from Szilágyi & Lakatos (2015) to enable analytic cross-moment closure without introducing unmeasured corner angles.

Hypothesis 3: Negligible Crystal Agglomeration and Breakage

Assumption: Crystals neither collide to form agglomerates nor shatter into fragments; particle count increases solely via secondary nucleation.

Engineering Justification: Valid for dilute-to-moderate slurry densities (solid volume fraction $V_c < 10\%$) operating below the critical shear fracture threshold.

Hypothesis 4: Secondary Contact Nucleation Dominance

Assumption: Primary homogeneous nucleation is zero; all newly born crystals originate from secondary contact nucleation:

$$B = k_S \varepsilon \mu_{1,1} \sigma^2.$$

Engineering Justification: The crystallizer is seeded at $T_s = 35.0^\circ\text{C}$ within the metastable zone, where supersaturation ($\sigma < 0.05$) is too low to trigger primary nucleation.

Hypothesis 5: Infinitesimal Nucleus Birth Size ($L_{10}, L_{20} \rightarrow 0$)

Assumption: Newly nucleated crystals enter the population at zero size.

Engineering Justification: Stable nuclei form at nanometer dimensions ($\sim 10^{-9} \text{ m}$). Relative to parent seeds ($\sim 10^{-4} \text{ m}$), their initial volume is negligible.

Hypothesis 6: Perfect Macroscopic Mixing (CSTR Hydrodynamics)

Assumption: Solute concentration $c(t)$, temperature $T(t)$, and crystal slurry density are spatially uniform throughout the 5-liter vessel.

Engineering Justification: The turbulent mixing timescale in a 5L agitated vessel ($\tau_{mix} \sim 1\text{--}3 \text{ s}$) is much faster than the cooling timescale ($\tau_{batch} \sim 12,000 \text{ s}$).

Hypothesis 7: Local Thermodynamic Phase Equilibrium at the Interface

Assumption: The solution saturation concentration $c_s(T)$ is instantaneously governed by the modified Apelblat equation at the current temperature.

Engineering Justification: Liquid-solid phase boundary relaxation occurs at molecular timescales ($\sim 10^{-8} \text{ s}$), far exceeding process cooling rates.

Hypothesis 8: Linear Quasi-Steady Temperature Control

Assumption: The jacket temperature tracks the linear setpoint trajectory exactly: $dT/dt = -cr$.

Engineering Justification: Standard assumption in simulation studies with tightly tuned slave jacket controllers.

Hypothesis 9: Exposure Bias in Recurrent Multi-Step Forecasting

Assumption: Black-box RNNs trained with Teacher Forcing compound errors during open-loop recursive rollouts; injecting Gaussian input noise ($\sigma_{\text{jitter}} = 0.012$) acts as contractive regularization that forces the GRU vector field to self-correct.

Machine Learning Justification: Validated by our experimental rollouts: noise regularization eliminated trajectory divergence across 80+ recursive steps.

Hypothesis 10: Inductive Bias Trade-off (Physics Regularization)

Assumption: Penalizing physical differential equation residuals ($\lambda_{\text{phys}} = 0.05$) sacrifices a negligible amount of single-variable curve-fitting accuracy to guarantee solute mass conservation and coupled facet growth.

Machine Learning Justification: While BB-RNN had a $4\mu\text{m}$ lower length error, it created 14x more unphysical mass. PI-RNN achieves 93% mass compliance and boosts Aspect Ratio R^2 from 0.11 to 0.72.

8. Professor Defense & Oral Exam Technical Cheat Sheet

Q1: "Why use the Method of Moments (QMOM) instead of discretizing the 2D PDE on a fine spatial grid?"

Discretizing a 2D PBE with length and width over a standard 100×100 mesh creates 10,000 coupled stiff ODEs. This incurs heavy computational cost (~15 minutes per batch) and causes numerical diffusion that artificially broadens the seed distribution. The Quadrature Method of Moments (QMOM) tracks the low-order statistical moments using 3 quadrature nodes, reducing the dynamics to 9 stiff ODEs solved in under 1.5 seconds with zero numerical diffusion while preserving the exact mean length, width, and mass balance.

Q2: "Why is the ODE system stiff, and why is Backward Differentiation Formulas (BDF) necessary?"

There is a massive separation of timescales. Secondary nucleation $B(t)$ and face growth respond to supersaturation within fractions of a second (10^{-1} s), whereas batch cooling and desaturation take 12,000 seconds (3.3 hours). Explicit integrators (like explicit Euler or RK45) undergo numerical instability unless the time step is $\ll 10^{-3}$ s. BDF is an implicit multistep algorithm designed specifically for stiff reaction and crystallization ODEs, ensuring numerical A-stability.

Q3: "Why did the Black-Box RNN fail on crystal aspect ratio ($R^2 = 0.11$), while the PI-RNN succeeded ($R^2 = 0.72$)?"

Over a 3-hour batch, the aspect ratio changes only from 1.78 to 2.25 (total variation $\Delta AR \approx 0.47$). The Black-Box RNN models length and width independently; slight uncoordinated errors between them cause the aspect ratio error to exceed the natural variance, destroying R^2 . The PI-RNN embeds the coupled growth kinetics G_1 and G_2 through supersaturation σ , ensuring face coupling and suppressing aspect ratio drift.

Q4: "BB-RNN had a slightly lower RMSE on length ($8\mu\text{m}$ vs $12\mu\text{m}$). Does that make it superior?"

No, this reflects the classic bias-variance trade-off in machine learning. Unconstrained black-box networks overfit single trajectories by violating the laws of physics—specifically by violating solute mass conservation by over 14x (residual: 0.184 vs 0.0128 kg/(m³·s)). The PI-RNN enforces the coupled mass balance and 2D kinetic equations; it sacrifices a negligible 4 microns of length error to guarantee that the multi-step trajectory strictly conserves mass (93% reduction in violation) and predicts crystal width within $2.0\mu\text{m}$.

Q5: "How was the dataset generated, and is it the same technique as the original PI-RNN paper?"

Yes, exactly. Following Zheng & Wu (2023), we sample the operational boundary space (seeding temperature, cooling rate, stirring power, seed mass), run our first-principles 2D QMOM simulation to produce continuous trajectories, slice them into moving sequence windows of length 15 (stride 2), and apply Z-score normalization strictly on the training partition without data leakage. The only difference is that Zheng & Wu evaluated a CSTR reaction system, while we adapted the PI-RNN to the 2D crystallization benchmark of Szilágyi & Lakatos (2015).

9. Complete Codebase Architecture & File Lookup

FILENAME	PRIMARY FUNCTION & MATHEMATICAL SCOPE
config.py	Process constants (Table 1), Apelblat coefficients (Eq 11), kinetic parameters, Table 2 quadrature seeds, and PI-RNN hyperparameters.
crystallizer_mom.py	First-principles 2D QMOM simulation engine using <code>scipy.integrate.solve_ivp(method='BDF')</code> for 9 stiff states.
reproduce_paper_results.py	Generates benchmark Figures 1, 3, 6, 8, 9, 10, and 11 from Szilágyi & Lakatos (2015).
dataset_generator.py	Simulates diverse batch trajectories across parameter envelopes; builds sliding sequences with z-score normalization.
pirnn_model.py	PyTorch 2-layer GRU with differentiable <code>PhysicsLossModule</code> enforcing mass balance and face growth residuals.
train_pirnn.py	Training pipeline for PI-RNN and Black-Box RNN with input noise jittering and convergence plotting.
compare_results.py	Closed-loop recursive rollout evaluator, metric table generator, and parity plot builder.