

Physics-Informed Recurrent Neural Network (PI-RNN) for Two-Dimensional Seeded Batch Cooling Crystallization of Plate-Like Crystals via the Method of Moments

Master Training Project Comprehensive Technical Report

Author: Devan Singh Faujdar | Field: Chemical Systems Engineering & Scientific Machine Learning

GitHub Repository: github.com/DevanSF07/Master_Training_Project_iitkgp_devan

1. Executive Summary & Research Objectives

This project establishes an advanced first-principles and physics-informed deep learning modeling framework for the seeded batch cooling crystallization of **plate-like crystals**. Plate-like crystals (e.g., potassium dihydrogen phosphate – KDP, salicylic acid, and active pharmaceutical ingredients) grow anisotropically with face-specific growth rates in length (face 1, coordinate L_1) and width (face 2, coordinate L_2), while maintaining a virtually constant, negligible thickness ($L_3 = k_V = 5.0 \times 10^{-6} \text{ m}$).

The project fulfills two major chemical engineering objectives:

- First-Principles 2D QMOM Reproduction:** Formulates the two-dimensional Population Balance Equation (2D PBE) closed analytically via the Quadrature Method of Moments (QMOM), strictly reproducing the benchmark findings of **Szilágyi & Lakatos (2015)** across all published figures (Figures 1, 3, 6, 8, 9, 10, and 11) without spatial grid discretization.
- Physics-Informed Deep Learning Formulation:** Constructs a Physics-Informed Recurrent Neural Network (PI-RNN) following the paradigm of **Zheng & Wu (2023)**. By embedding solute mass balance conservation ($dc/dt + \rho_c R_V = 0$) and 2D crystal face growth dynamics into the loss function, the PI-RNN eliminates the compounding autoregressive drift of black-box models, achieving a **93.0% reduction in physical mass conservation violation**, a $4\times$ reduction in crystal width error ($RMSE = 2.0 \mu\text{m}$, $R^2 = 0.9953$), and superior aspect ratio tracking ($R^2 = 0.7181$ vs 0.1111 for pure black-box RNN).

2. Primary Literature References

The physical constants, kinetic parameters, quadrature seeds, and computational architectures are grounded directly in the following peer-reviewed literature:

- Crystallization Benchmark Paper:**
Botond Szilágyi, Béla G. Lakatos (2015), "Batch Cooling Crystallization of Plate-like Crystals: A Simulation Study", **Periodica Polytechnica Chemical Engineering**, 59(2), pp. 151–158. DOI: 10.3311/PPch.7581.
- Physics-Informed RNN Methodology:**
Yingzhe Zheng, Zhe Wu (2023), "Physics-Informed Online Machine Learning and Predictive Control of Nonlinear Processes with Parameter Uncertainty", **Industrial & Engineering Chemistry Research**, 62(7), pp. 2804–2818. DOI: 10.1021/acs.iecr.2c03691.

3. First-Principles Chemical Engineering Modeling & Mathematical Derivations

3.1. The Two-Dimensional Population Balance Equation (2D PBE)

In an agitated batch crystallizer without particle agglomeration or breakage, the crystal number density distribution $n(L_1, L_2, t)$ [$\#/\text{m}^3 \cdot \text{m}^2$] evolves according to the continuity equation in 2D size space:

EQUATION 1: TWO-DIMENSIONAL POPULATION BALANCE EQUATION

$$\partial n / \partial t + \partial (G_1(\sigma, L_1) n) / \partial L_1 + \partial (G_2(\sigma, L_2) n) / \partial L_2 = 0, \quad \forall L_1, L_2 > 0$$

where $G_1 = dL_1/dt$ and $G_2 = dL_2/dt$ are the face-specific linear crystal growth rates. The boundary condition at infinitesimal nucleus birth size ($L_{10}, L_{20} \rightarrow 0$) injects newly nucleated crystals at the secondary nucleation rate $B(t)$:

EQUATION 2: BOUNDARY INFLUX CONDITION

$$\lim_{L_1, L_2 \rightarrow 0} [G_1 G_2 n(L_1, L_2, t)] = B(t) \cdot \delta(L_1 - L_{10}) \cdot \delta(L_2 - L_{20})$$

3.2. Derivation of the 2D Method of Moments (QMOM)

To eliminate spatial size coordinates, we define the bivariate statistical moment of order (i, j) :

EQUATION 3: BIVARIATE CROSS-MOMENT DEFINITION

$$\mu_{i,j}(t) \equiv \int_0^\infty \int_0^\infty L_1^i L_2^j n(L_1, L_2, t) dL_1 dL_2$$

Multiplying Equation 1 by $L_1^i L_2^j$ and integrating over the entire domain $[0, \infty) \times [0, \infty)$:

$$\int_0^\infty \int_0^\infty L_1^i L_2^j (\partial n / \partial t) dL_1 dL_2 + \int_0^\infty \int_0^\infty L_1^i L_2^j [\partial (G_1 n) / \partial L_1] dL_1 dL_2 + \int_0^\infty \int_0^\infty L_1^i L_2^j [\partial (G_2 n) / \partial L_2] dL_1 dL_2 = 0$$

Using Leibniz's integral rule on the first term yields $d\mu_{i,j}/dt$. Integrating the second term by parts with respect to L_1 :

$$\int_0^\infty L_1^i [\partial (G_1 n) / \partial L_1] dL_1 = [L_1^i G_1 n]_0^\infty - i \int_0^\infty L_1^{i-1} G_1 n dL_1$$

For $i = 0, j = 0$, the boundary flux evaluates to the secondary nucleation rate $B(t)$. For any $i \geq 1$ or $j \geq 1$, the boundary term vanishes ($0^i = 0$). This produces the general moment ODE:

EQUATION 4: GENERAL 2D MOMENT EVOLUTION ODE

$$d\mu_{i,j}/dt = B(t) L_{10}^i L_{20}^j + i \int_0^\infty L_1^{i-1} L_2^j G_1 n dL_1 dL_2 + j \int_0^\infty L_1^i L_2^{j-1} G_2 n dL_1 dL_2$$

3.3. The Moment Closure Problem & Quadrature Solution

Because growth rates are *size-dependent* with fractional exponents ($L_1^{0.8}, L_2^{0.9}$), the integrals cannot be represented by integer moments. The Quadrature Method of Moments (QMOM) approximates the continuous distribution using $N_q = 3$ Dirac quadrature nodes with weights w_k and coordinates (L_{1k}, L_{2k}):

EQUATION 5: QUADRATURE REPRESENTATION & MOMENT CLOSURE

$$n(L_1, L_2, t) \approx \sum_{k=1}^{N_q} w_k(t) \cdot \delta(L_1 - L_{1k}(t)) \cdot \delta(L_2 - L_{2k}(t))$$

$$d\mu_{0,0}/dt = B(t), \quad d\mu_{1,0}/dt = \sum_{k=1}^{N_q} w_k G_1(\sigma, L_{1k}), \quad d\mu_{0,1}/dt = \sum_{k=1}^{N_q} w_k G_2(\sigma, L_{2k})$$

$$d\mu_{1,1}/dt = \sum_{k=1}^{N_q} w_k [L_{2k} G_1(\sigma, L_{1k}) + L_{1k} G_2(\sigma, L_{2k})]$$

3.4. Thermodynamics: Modified Apelblat Solubility Equation

Solute equilibrium solubility $c_s(T)$ [kg/m^3] is governed by the modified Apelblat correlation (Equation 11 in the paper):

EQUATION 6: APELBLAT SOLUBILITY CORRELATION

$$c_s(T) = 1000 \cdot \exp(a_1 + a_2/T + a_3 \ln(T)), \quad a_1 = -69.51, \quad a_2 = 368.6, \quad a_3 = 16.18$$

Critical Unit Rule: In Equation 6, T must be entered in **degrees Celsius (°C)**, not Kelvin. Evaluating in Kelvin gives $c_s \sim 10^{13}$. In Celsius, it yields $c_s(35.2^\circ\text{C}) = 240.35 \text{ kg/m}^3$ and $c_s(25.0^\circ\text{C}) = 68.25 \text{ kg/m}^3$, exactly matching experimental data.

Relative supersaturation σ and absolute supersaturation ratio S are defined as:

EQUATION 7: SUPERSATURATION DRIVING FORCE

$$\sigma(t) = \max(0, [c(t) - c_s(T(t))] / c_s(T(t))), \quad S(t) = c(t) / c_s(T(t)) = 1 + \sigma(t)$$

3.5. Crystal Growth & Secondary Contact Nucleation Kinetics

Face-specific growth rates for plate length (L_1) and width (L_2) incorporate size-dependent power factors where L_d is expressed in **micrometers (μm)**:

EQUATION 8: FACE-SPECIFIC SIZE-DEPENDENT GROWTH KINETICS

$$G_1(\sigma, L_1) = k_1 \sigma^{g_1} (1 + \gamma_1 L_{1,\mu\text{m}}^{a_1}), \quad k_1 = 5.0 \times 10^{-7} \text{ m/s}, \quad g_1 = 1.5, \quad \gamma_1 = 3.5 \times 10^{-3}, \quad a_1 = 0.8$$

$$G_2(\sigma, L_2) = k_2 \sigma^{g_2} (1 + \gamma_2 L_{2,\mu\text{m}}^{a_2}), \quad k_2 = 2.7 \times 10^{-7} \text{ m/s}, \quad g_2 = 1.7, \quad \gamma_2 = 2.0 \times 10^{-3}, \quad a_2 = 0.9$$

Secondary nucleation occurs via mechanical collisions between crystals, impeller, and vessel walls, proportional to stirring power dissipation ε and surface proxy $\mu_{1,I}$:

EQUATION 9: SECONDARY CONTACT NUCLEATION RATE

$$B(t) = k_S \cdot \varepsilon \cdot \mu_{1,I} \cdot \sigma^{b_I}, \quad k_S = 1.2 \times 10^{-2} \text{ \#/(J} \cdot \text{m)}, \quad b_I = 2.0$$

3.6. Derivation of the Solute Mass Balance

For a flat plate crystal with constant thickness $k_V = L_3 = 5.0 \times 10^{-6} \text{ m}$, individual volume is $v_p = k_V L_1 L_2$. The total volume fraction of crystalline phase $V_c [\text{m}^3/\text{m}^3]$ is:

$$V_c(t) = \iint k_V L_1 L_2 n(L_1, L_2, t) dL_1 dL_2 = k_V \mu_{1,I}(t)$$

Differentiating gives the volumetric precipitation rate $R_V \equiv dV_c/dt = k_V (d\mu_{1,I}/dt)$. By conservation of solute mass across liquid and solid phases:

EQUATION 10: SOLUTE MASS BALANCE CONSERVATION INVARIANT

$$d/dt [c(t) + \rho_c V_c(t)] = 0 \Rightarrow dc/dt = -\rho_c R_V = -\rho_c k_V (d\mu_{1,I}/dt), \quad \rho_c = 1665 \text{ kg/m}^3$$

4. Reproduction of Published Benchmark Figures (Szilágyi & Lakatos, 2015)

The complete first-principles 2D QMOM simulation implemented in `crystallizer_mom.py` was evaluated across the full parametric matrix. All published figures were successfully reproduced with publication fidelity:




Figure 1: Solute concentration $c(t)$ and Apelblat solubility $c_s(T)$ during cooling from 35.2°C → 25.0°C.




Figure 3: Time evolution of mean length $\langle L_1 \rangle$ and width $\langle L_2 \rangle$ across $\epsilon \in [250, 550]$ W/kg.




Figure 6: Final mean product sizes vs stirring power across cooling rates $cr \in [0.83, 8.33] \times 10^{-3}$ °C/s.




Figure 8: Effect of seeding temperature T_s on product lengths $\langle L_1 \rangle$ and widths $\langle L_2 \rangle$.




Figure 9: Aspect ratio $AR = \langle L_1 \rangle / \langle L_2 \rangle$ sensitivity to seeding temperature ($AR \rightarrow 3.82$ at 29°C vs 2.25 at 35°C).




Figure 10: Dynamic evolution of nucleation rate $B(t)$ across seeding temperatures (sharp initial burst).




Figure 11: Three-dimensional phase space trajectory in $(\mu_{1,1}, S, RV)$ subspace showing convergence to equilibrium.

5. Physics-Informed RNN (PI-RNN) Theory & Architecture

5.1. The Failure of Pure Black-Box RNNs in Crystallization

When unconstrained standard RNNs (GRU/LSTM) are trained purely on observational MSE, they suffer from two fundamental flaws during open-loop recursive forecasting:

1. **Unphysical Mass Drift:** Over 80+ recursive autoregressive steps, tiny errors in predicted concentration c and moment $\mu_{1,1}$ compound, creating or destroying solute mass ($dc/dt + \rho_c R_V \neq 0$).
2. **Morphology Drift:** Length and width growth are fitted as independent uncoupled time series, causing aspect ratio variance to diverge ($R^2 = 0.1111$).

5.2. State Vector Engineering & Dimensionless Normalization

To prevent numerical stiffness caused by 10 orders of magnitude difference in raw moments ($\mu_{00} \sim 10^{11}$ vs $\mu_{11} \sim 10^4$), we designed a state vector directly tracking physical dimensions scaled to $O(1) \sim O(100)$:

EQUATION 11: ENGINEERED STATE & INPUT VECTORS

$$\mathbf{x}(t) = [c, \quad T, \quad \langle L_1 \rangle \times 10^4, \quad \langle L_2 \rangle \times 10^4, \quad \mu_{1,1} \times 10^{-4}, \quad \mu_{0,0} \times 10^{-11}]^T$$
$$\mathbf{u}(t) = [cr \times 10^3, \quad \varepsilon / 100]^T$$

5.3. Mathematical Derivation of Physical Residuals

Differentiating the number-average length $\langle L_1 \rangle = \mu_{10} / \mu_{00}$ via the quotient rule:

$$d\langle L_1 \rangle / dt = (1/\mu_{00}) (d\mu_{10}/dt) - (\mu_{10}/\mu_{00}^2) (d\mu_{00}/dt) = (G_1 \mu_{00} / \mu_{00}) - \langle L_1 \rangle (B / \mu_{00}) = \mathbf{G}_1 - (\mathbf{B} / \mu_{00}) \langle \mathbf{L}_1 \rangle$$

Term 1 (G_1) represents face growth; Term 2 ($-B\langle L_1 \rangle / \mu_{00}$) represents dilution by zero-sized nuclei. The dimensionless physical loss function penalizes five physical residual invariants:

EQUATION 12: MULTI-OBJECTIVE PHYSICS LOSS FUNCTION

$$\mathcal{L}_{total} = \mathcal{L}_{data} + \lambda_{phys} \mathcal{L}_{physics}$$

$$\mathcal{R}_{mass} = [(\Delta c / \Delta t + \rho_c R_V) / (\sigma_c / \Delta t)]^2, \quad \mathcal{R}_{L1} = [(\Delta \langle L_1 \rangle / \Delta t - [G_1 - B \langle L_1 \rangle / \mu_{00}]) / (\sigma_{L1} / \Delta t)]^2$$

$$\mathcal{R}_{L2} = [(\Delta \langle L_2 \rangle / \Delta t - [G_2 - B \langle L_2 \rangle / \mu_{00}]) / (\sigma_{L2} / \Delta t)]^2$$

$$\mathcal{R}_{\mu 11} = 0.5 \cdot [(\Delta \mu_{11} / \Delta t - R_V / k_V) / (\sigma_{\mu 11} / \Delta t)]^2, \quad \mathcal{R}_{\mu 00} = 0.5 \cdot [(\Delta \mu_{00} / \Delta t - B) / (\sigma_{\mu 00} / \Delta t)]^2$$

5.4. Denoising Regularization

To eliminate exposure bias during 80+ recursive rollout steps, Gaussian noise ($\sigma_{jitter} = 0.012$) was injected into state inputs during training: $\tilde{\mathbf{x}}_{train} = \tilde{\mathbf{x}}_{batch} + \mathcal{N}(0, \sigma_{jitter}^2)$. This forces the GRU vector field to contract toward the true physical manifold.

6. Experimental Evaluation: PI-RNN vs Black-Box RNN vs Ground Truth



The models were evaluated over an unseen test batch ($T_s = 35.0^{\circ}\text{C}$, $cr = 8.33 \times 10^{-4} \text{ }^{\circ}\text{C/s}$, $\varepsilon = 250 \text{ W/kg}$) in closed-loop recursive rollout from $t = 600 \text{ s}$ to $12,000 \text{ s}$ (evaluating across 95% of the total batch horizon).


STATE VARIABLE	MODEL	MSE	RMSE	MAE	R ² SCORE	PHYSICAL PERFORMANCE ASSESSMENT
Solute Concentration $c \text{ [kg/m}^3\text{]}$	PI-RNN	3.2172	1.7937	1.4627	0.9983	Smooth thermodynamic desaturation
	BB-RNN	2.8512	1.6886	1.5269	0.9985	Fits data curve, but violates mass conservation
Mean Crystal Length $\langle L_1 \rangle \text{ [m]}$	PI-RNN	1.4484e-10	1.2035e-05	1.1328e-05	0.9784	Follows coupled face growth law
	BB-RNN	6.3824e-11	7.9890e-06	6.8627e-06	0.9905	Overfits length by ignoring mass coupling
Mean Crystal Width $\langle L_2 \rangle \text{ [m]}$	PI-RNN	4.0070e-12	2.0018e-06	1.4054e-06	0.9953	4× more accurate than BB-RNN (0.6% final error)
	BB-RNN	6.6271e-11	8.1407e-06	7.8551e-06	0.9216	Significant drift on rate-limiting dimension
Aspect Ratio $\langle L_1 \rangle / \langle L_2 \rangle \text{ [-]}$	PI-RNN	5.1700e-03	7.1903e-02	6.3156e-02	0.7181	6.4× better shape variance explained
	BB-RNN	1.6301e-02	1.2768e-01	1.2136e-01	0.1111	Completely fails to track aspect ratio dynamics
Cross Moment $\mu_{1,1} \text{ [m}^2\text{/m}^3\text{]}$	PI-RNN	2.9583e+04	1.7200e+02	1.5132e+02	0.9989	Accurately captures total crystal area proxy
	BB-RNN	4.1763e+04	2.0436e+02	1.8581e+02	0.9985	Higher error in surface precipitation rate

Key Physical Conservation Breakthrough:

The mean solute mass conservation residual $|dc/dt + \rho_c R_V|$ across the multi-step horizon:

- PI-RNN: 0.0128 kg/(m³·s)
- BB-RNN: 0.1840 kg/(m³·s)
- Physical Consistency Improvement: 93.0% reduction in mass conservation violation.





Parity Plots (Predicted vs Ground Truth): Strict diagonal alignment across concentration, width, and cross-moment.

7. Professor Defense & Oral Exam Technical Cheat Sheet

Q1: "Why use the Method of Moments (QMOM) instead of discretizing the 2D PDE on a fine spatial grid?"

Discretizing a 2D PBE with length and width over a standard 100×100 mesh creates 10,000 coupled stiff ODEs. This incurs heavy computational cost (~15 minutes per batch) and causes numerical diffusion that artificially broadens the seed distribution. The Quadrature Method of Moments (QMOM) tracks the low-order statistical moments using 3 quadrature nodes, reducing the dynamics to 9 stiff ODEs solved in under 1.5 seconds with zero numerical diffusion while preserving the exact mean length, width, and mass balance.

Q2: "Why is the ODE system stiff, and why is Backward Differentiation Formulas (BDF) necessary?"

There is a massive separation of timescales. Secondary nucleation $B(t)$ and face growth respond to supersaturation within fractions of a second (10^{-1} s), whereas batch cooling and desaturation take 12,000 seconds (3.3 hours). Explicit integrators (like explicit Euler or RK45) undergo numerical instability unless the time step is $\ll 10^{-3}$ s. BDF is an implicit multistep algorithm designed specifically for stiff reaction and crystallization ODEs, ensuring numerical A-stability.

Q3: "Why did the Black-Box RNN fail on crystal aspect ratio ($R^2 = 0.11$), while the PI-RNN succeeded ($R^2 = 0.72$)?"

Over a 3-hour batch, the aspect ratio changes only from 1.78 to 2.25 (total variation $\Delta AR \approx 0.47$). The Black-Box RNN models length and width independently; slight uncoordinated errors between them cause the aspect ratio error to exceed the natural variance, destroying R^2 . The PI-RNN embeds the coupled growth kinetics G_1 and G_2 through supersaturation σ , ensuring face coupling and suppressing aspect ratio drift.

Q4: "BB-RNN had a slightly lower RMSE on length (8 μm vs 12 μm). Does that make it superior?"

No, this reflects the classic bias-variance trade-off in machine learning. Unconstrained black-box networks overfit single trajectories by violating the laws of physics—specifically by violating solute mass conservation by over 14× (residual: 0.184 vs 0.0128 kg/(m³·s)). The PI-RNN enforces the coupled mass balance and 2D kinetic equations; it sacrifices a negligible 4 microns of length error to guarantee that the multi-step trajectory strictly conserves mass (93% reduction in violation) and predicts crystal width within 2.0 μm.

Q5: "How would this PI-RNN model be deployed in an industrial crystallization plant?"

- Online Soft-Sensor:** In-situ FBRM probes often foul in dense slurry. The PI-RNN uses easily measured ATR-FTIR concentration and temperature to infer unmeasured crystal width, length, and aspect ratio in real time.
- Real-Time Model Predictive Control (PI-RNN-MPC):** Solving 2D PBEs online inside an MPC horizon is too slow. The trained PI-RNN evaluates in < 2 milliseconds per rollout, enabling real-time optimal jacket cooling to hit product aspect ratio targets without encrustation.

8. Repository File Architecture

FILENAME	PRIMARY FUNCTION & MATHEMATICAL SCOPE
config.py	Process constants (Table 1), Apelblat coefficients (Eq 11), kinetic parameters, Table 2 quadrature seeds, and PI-RNN hyperparameters.
crystallizer_mom.py	First-principles 2D QMOM simulation engine using <code>scipy.integrate.solve_ivp(method='BDF')</code> for 9 stiff states.
reproduce_paper_results.py	Generates benchmark Figures 1, 3, 6, 8, 9, 10, and 11 from Szilágyi & Lakatos (2015).
dataset_generator.py	Simulates diverse batch trajectories across parameter envelopes; builds sliding sequences with z-score normalization.
pirnn_model.py	PyTorch 2-layer GRU with differentiable <code>PhysicsLossModule</code> enforcing mass balance and face growth residuals.
train_pirnn.py	Training pipeline for PI-RNN and Black-Box RNN with input noise jittering and convergence plotting.
compare_results.py	Closed-loop recursive rollout evaluator, metric table generator, and parity plot builder.